A Quadratic Assignment Problem Approach to Graph Matching: Applications in Statistical Connectomics

Joshua T. Vogelstein, John C. Conroy, Lou Podrazik, Steve Kratzer, R. Jacob Vogelstein, and Carey E. Priebe

Abstract—It is becoming increasingly popular to represent myriad and diverse data sets as graphs. When the labels of vertices of these graphs are unavailable, graph matching (GM)—the process of determining whether the graphs are isomorphic to one another— is a computationally daunting problem. This work presents an inexact strategy for GM. Specifically, we relax of the feasible region to its convex hull and then apply a well known and efficient nonlinear programming algorithm, Frank-Wolfe, to the relaxed objective function. Though this relaxation is convex, the point around which the local approximation is made determines the apex. We therefore consider a number of initializations based on the geometry of the convex hull. Multiple restarts of this algorithm leads to performance that exceeds the previous state-of-the-art in *all* of 16 benchmark tests. Moreover, this approach is fast, scaling cubically with the number of vertices, requiring only a few minutes on standard modern laptops for graphs with up to a few hundred vertices. We then applied this approach to a brain-graph in which vertices represent neurons in a small nematode brain (the *Caenorhabditis elegans* worm), and edges represent either chemical or electrical synapses. For every chemical connectome and several electrical connectomes, this approach found the optimal solution. Although this strategy already natively operates on unweighted and weighted graphs, directed and undirected graphs, we propose a number of possible extensions, and make all code available.

Index Terms—statistical inference, graph theory, network theory, structural pattern recognition, connectome.

✦ ─────────────

1 INTRODUCTION

A graph matching (GM) algorithm is any algorithm who's goal is to "align" any set of $n \geq 2$ graphs such that each vertex v in one graph can be "assigned" to its corresponding vertex in the other graphs. Perhaps due to its complex computational properties (it is NP-hard [1]), GM has received widespread attention in both the mathematical graph theory and computer science communities [2]. Moreover, the potential span of applications of graph matching algorithms is vast, ranging from neural coding [3] to machine vision [4].

Our motivation is the bourgeoning field called connectomics: the study of brain-graphs. In brain-graphs, vertices represent (collections of) neurons and edges represent either functional dependencies or structural connections [5]. In some scenarios vertices are labeled. For example, when vertices represent single neurons in invertebrates [6], or macro-anatomical gyral regions in vertebrates [7], [8]. However, in other scenarios, even whether vertices can be labeled is questionable. For example, if one desired to compare brain-(sub)graphs from parts of brains across species or mammalian organisms, there is no known vertex assignment. In these scenarios

GM might be an important element of any statistical analysis of these brain-graphs [9].

We therefore propose a novel inexact graph matching algorithm. The intuition is quite simple: GM is computationally difficult because the set of feasible solutions is (i) non-differentiable and (ii) multimodal (and large). A common approach to approximating difficult nonlinear programming problems is to relax the constraints on the feasible region. By relaxing the non-differentiable constraint, any gradient based algorithm may be applied to the problem [10]. Unfortunately, the multimodality of the solution space implies that the initialization will, in general, be important. Multiple "principled" restarts can potentially facilitate an efficient stochastic search strategy.

This manuscript describes an algorithm that approximately solves graph matching in cubic time (with very small leading constants). Via numerical experiments, we demonstrate that this approach outperforms several state-of-the-art algorithms on all tests in a standard benchmark library [11]. Moreover, it runs in cubic time with very small leading constants, making it both fast and effective. We then test this approach on a brain-graph matching problem: matching the brain-graphs of a small nematode with 302 vertices. We are able to find the optimal solution after 3 restarts for each randomly permuted example. We are therefore optimistic that this algorithm will be useful for the large graphs promised to arise due to various ongoing connectome projects [12], [13].

- *J.T. Vogelstein and C.E. Priebe are with the Department of Applied Mathematics and Statistics, Johns Hopkins University, Baltimore, MD 21218. E-mail: {joshuav,cep}@jhu.edu*
- *R.J. Vogelstein is with the Johns Hopkins University Applied Physics Laboratory, Laurel, MD, 20723.*

2 METHODS

2.1 Preliminaries

A labeled graph $G = (\mathcal{V}, \mathcal{E})$ consists of a vertex set, $\mathcal{V} = [n] = \{1, \ldots, n\}$, where n is number of vertices and an edge set $\mathcal{E} \subseteq \binom{[n]}{2}$. Let A be the adjacency matrix representations of graph such that $A_{uv} = 1$ if $u \sim v \in \mathcal{E}$ and $A_{uv} = 0$ otherwise. Let $Q \in \mathcal{Q}$ be a permutation matrix (a matrix with a single one in each row and column). Given a pair of adjacency matrices, A and B, to graph match A with B is to find a permutation matrix Q such that $QAQ^\mathsf{T} = B$. In this work, we propose a novel inexact graph matching algorithm, essentially a Frank-Wolfe algorithm with multiple restarts. We demonstrate the efficacy of this algorithm over the previous-state-of-the-art on a reference library of benchmarks.

2.2 A QAP Approach to GM

Graph matching can be cast as a quadratic assignment problem [14]:

$$Q_{AB} = \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} \left\| QAQ^\mathsf{T} - B \right\|_F^2, \qquad (1)$$

where the permutation matrix Q_{AB} induces a labeling of the vertices of A onto those of B. A bit of linear algebra simplifies Eq. (1):

$$\begin{aligned}
\underset{Q \in \mathcal{Q}}{\operatorname{argmin}} &\left\| QAQ^\mathsf{T} - B \right\|_F^2 \\
&= \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} -tr(B^\mathsf{T} QAQ^\mathsf{T}) - tr(QAQ^\mathsf{T} B),
\end{aligned}$$

which follows from the definition of the Frobenius norm and canceling appropriately. Note that the objective function on the right-hand-side is equivalent to the standard representation of the quadratic assignment problem (QAP) [2]:

$$\hat{\sigma} = \underset{\sigma}{\operatorname{argmin}}\, a_{\sigma(u),\sigma(v)} b_{uv} = \underset{q \in \mathcal{Q}}{\operatorname{argmin}}\, q_{uv} a_{uv}, q_{vu} b_{uv} \qquad (2)$$

where σ is a permutation function, $\sigma : [n] \mapsto [n]$. Unfortunately, Eq. (1) is an NP-complete problem [1]. The primary difficulty in solving Eq. (1) is the discrete non-convex constraint set. Thus, it is natural to consider an approximation with the constraints relaxed. Since the convex hull of permutation matrices is the set of doubly stochastic matrices (matrices whose rows and columns all sum to one), we define the approximate quadratic assignment problem:

$$\widetilde{Q}_{AB} = \underset{Q \in \mathcal{D}}{\operatorname{argmin}} \left\| QAQ^\mathsf{T} - B \right\|_F^2, \qquad (3)$$

where $\mathcal{D}$ is the set of doubly stochastic matrices. When the permutation matrix constraint is relaxed, the equivalence relation shown in Eq. (2) no longer holds. Nonetheless, we proceed by attempting to solve:

$$\hat{\widetilde{Q}}_{AB} \approx \underset{Q \in \mathcal{D}}{\operatorname{argmin}} -tr(B^\mathsf{T} QAQ^\mathsf{T}) - tr(QAQ^\mathsf{T} B), \qquad (4)$$

considering it an auxiliary function for which we can compute gradients and ascend a likelihood, unlike the permutation constrained case.

The Frank-Wolfe (FW) algorithm is a successive linear programming algorithm originally designed for solving quadratic problems with linear (equality and/or inequality) constraints [15]. It later became used more more generally for nonlinear programming problems [16]. Let $f(Q) = -tr(B^\mathsf{T} QAQ^\mathsf{T}) - tr(QAQ^\mathsf{T} B)$ be our objective function. Unfortunately, this objective function is not necessarily positive definite. This is clear upon computing the Hessian of f with respect to Q:

$$\nabla_Q^2 = B \otimes A + B^\mathsf{T} \otimes A^\mathsf{T}, \qquad (5)$$

where $\otimes$ indicates the Kronecker product. This means that it will potentially be multi-modal, making initialization potentially important. We therefore develop a nonlinear programing algorithm for approximately solving Eq. (1) using multiple restarts. Below we provide details and explanation of each step.

Step 0: Choose an initial estimate While any doubly stochastic matrix would be a feasible initial point, two choices seem natural: (i) the "flat doubly stochastic matrix," $J = \mathbf{1}^\mathsf{T}\mathbf{1}/n$, which is the middle of the feasible region, and (ii) the identity matrix, which is a permutation matrix. Therefore, if we run QAP once, we always start with one of those two. If we use multiple restarts, each initial point is "near" the flat matrix. Specifically, we sample J', a random doubly stochastic matrix using 10 iterations of Sinkhorn balancing [17], and let $Q^{(0)} = (J + J')/2$. Given this initial estimate, we iterate the following five steps until convergence.

Step 1: Compute the gradient The gradient of f with respect to Q is given by:

$$\nabla_Q^{(j)} = \partial f / \partial Q^{(j)} = -AQ^{(j)} B^\mathsf{T} - A^\mathsf{T} Q^{(j)} B. \qquad (6)$$

Step 2: Find the closest doubly stochastic matrix Instead of directly descending this gradient, we search for the direction of the doubly stochastic matrix closest to this gradient. Noting that that direction may be computed by the dot-product operator, we have:

$$L^{(j)} = \underset{L \in \mathcal{D}}{\operatorname{argmin}} \langle L, \nabla_Q^{(j)} \rangle. \qquad (7)$$

Eq. (7) can be solved as a Linear Assignment Problem (LAP). More specifically, a LAP can be written as:

$$L_{AB} = \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} \left\| QA - B \right\|_F^2, \qquad (8)$$

which, when B is the identity matrix I, the above can be simplified:

$$\begin{aligned}
L_{AI} &= \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} \left\| QA - I \right\|_F^2 = \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} (QA - I)^\mathsf{T}(QA - I) \\
&= \underset{Q \in \mathcal{Q}}{\operatorname{argmin}}\, tr(A^\mathsf{T} Q^\mathsf{T} QA) - tr(2QA) - tr(II) \\
&= \underset{Q \in \mathcal{Q}}{\operatorname{argmin}} -\langle Q, A \rangle, \qquad (9)
\end{aligned}$$

In other words, letting $B = I$, the projection of a matrix onto its nearest doubly stochastic matrix is a LAP problem. While Eq. (9) cannot be solved directly, as above, we can relax the permutation matrix constraint to the doubly stochastic matrix constraint:

$$\widetilde{L}_{AI} = \operatorname*{argmin}_{Q \in \mathcal{D}} -\langle Q, A \rangle. \tag{10}$$

Since the permutation matrices are the vertices of the set of doubly stochastic matrices, finding the minimum of Eq. (10) is guaranteed to yield a permutation matrix (as minima are necessarily at the vertices). Thus, letting $A = \nabla_Q^{(j)}$, solving Eq. (10)—which is a linear problem with linear and non-negative constraints—is equivalent to solving Eq. (7). Fortunately, the Hungarian algorithm solves any LAP in $\mathcal{O}(n^3)$ [14], thus this projection is relatively computationally efficient.[1]

Step 3: Update the direction Given $L^{(j)}$, the new direction is given by:

$$d^{(j)} = L^{(j)} - Q^{(j)}. \tag{11}$$

Step 4: Line search Given this direction, one can then perform a line search to find the doubly stochastic matrix that minimizes the objective function along that direction:

$$\alpha^{(j)} = \operatorname*{argmin}_{\alpha \in [0,1]} f(Q^{(j)} + \alpha^{(j)} d^{(j)}). \tag{12}$$

This can be performed exactly, because f is a quadratic function.

Step 5: Update Q Finally, the new estimated doubly stochastic matrix is given by:

$$Q^{(j+1)} = Q^{(j)} + \alpha^{(j)} d^{(j)}. \tag{13}$$

Stopping criteria Steps 1–5 are iterated until convergence, computational budget limits, or some other stopping criterion is met. These 5 steps collective comprise the FW algorithm for solving Eq. (4). Note that while $Q^{(j)}$ will generally not be a permutation matrix, we do not project $Q^{(j)}$ back onto the set of permutation matrices between each iteration, as that projection requires $\mathcal{O}(n^3)$ time.

Projecting onto the set of permutation matrices Let $Q^{(J+1)}$ be the doubly stochastic matrix resulting from the final iteration. We project $Q^{(J+1)}$ onto the set of permutation matrices:

$$\hat{Q} = \operatorname*{argmin}_{Q \in \mathcal{Q}} \langle Q^{(J+1)}, Q \rangle, \tag{14}$$

which is a LAP, and yields an approximate solution to the original QAP. Let FW appended with a projection onto the permutation matrices be denoted by `QAP`.

Multiple restarts We refer to multiple re-starts of `QAP` with subscripts, that is, the performance of `QAP`$_m$ is the best result of m pseudo-random re-starts of `QAP`. We set

---

1. More efficient algorithms are available for certain special cases, that is, whenever the matrix-vector multiplication operation is fast (for example, when both A and B are sparse).

m based on computational budget, lacking criteria for assessing whether we have obtained the global solution. Note that `QAP` natively operates on matrices, which could correspond to either weighted or unweighted graphs.

3 NUMERICAL RESULTS

3.1 QAP benchmarks

We first compare the performance of `QAP`$_m$ with recent state-of-the-art approaches on the QAP benchmark library [11]. Specifically, [18] reported improved performance in all but two cases, in which the QPB method of Cremers et al. [19] achieved a lower minimum. We compare `QAP`$_m$ with the previous best performing algorithm. In *all* cases, `QAP`$_3$ outperforms the previous best result, often by orders of magnitude in terms of relative error. In three cases, `QAP`$_{100}$ achieves the absolute minimum. In 12 out of 16 cases, 75%, the simple `QAP`$_1$ algorithm outperforms the others (starting with the flat doubly stochastic matrix). See Figure 1 for quantitative comparisons.


Fig. 1. `QAP`$_3$ outperforms PSOA on all 16 benchmark graph matching problems. Moreover, `QAP`$_1$ outperforms PSOA on 12 of 16 tests. For 3 of 16 tests, `QAP`$_{100}$ achieves the minimum (none of the other algorithms ever find the absolute minimum), as indicated by a black dot. Let f_* be the minimum and $\hat{f}_x$ be the minimum achieved by algorithm x. Error is $f_*/\hat{f}_x - 1$.

3.2 LAP vs. QAP

Much like the QAP objective function from Eq. (1) can be simplified to Eq. (2), the LAP objective function can be similarly simplified:

$$Q_{LAP} = \operatorname*{argmin}_{Q \in \mathcal{Q}} \|QA - B\|_F^2 = \operatorname*{argmin}_{Q \in \mathcal{Q}} tr(QAB^\mathsf{T}). \tag{15}$$

The gradient of the argument on the right-hand-side of the above equation is $2AB^\mathsf{T}$. Comparing this gradient to that of the QAP objective function—Eq. (6)—one can see that when (i) $Q^{(j)}$ is the identity matrix and (ii) both A and B are symmetric (for example, for undirected graphs), the two gradients are identical. Thus, if QAP is initialized at the identity matrix and the graphs are

undirected, the first permutation matrix—the output of Step 2—is identical to $\hat{Q}_{LAP}$; although the line search will make $Q^{(1)} \neq Q_{LAP}$, in general.

3.3 Algorithm Complexity and leading constants

Both GM and its closely related counterpart, graph isomorphism (GI), are computationally difficult. There exist no known algorithms for which worst case behavior is polynomial [?]. While GM is known to be NP-hard, it remains unclear whether GI is in P, NP, or its own intermediate complexity class, NP-isomorphism (or isomorphism-complete). Yet, for large classes of GI and GM problems linear or polynomial time algorithms are available [?]. Moreover, at worst, it is clear that GI is only "moderately exponential," for example, $\mathcal{O}(exp\{n^{1/2+o(1)}\})$ [?]. Unfortunately, even when linear or polynomial time GM or GI algorithms are available for special cases of graphs, the constants are typically unbearably large. For example, if all graphs have degree less than k, there is a linear time algorithm for GI. However, the hidden constant in this algorithm is $512k^3$! [?]. We therefore decided to determine the average complexity of our algorithm *and* the leading constants. Figure 2 suggests that our algorithm is not just cubic in time, but also has very small leading constants ($\approx 10^{-6}$), making using this algorithm feasible for even reasonably large graphs.
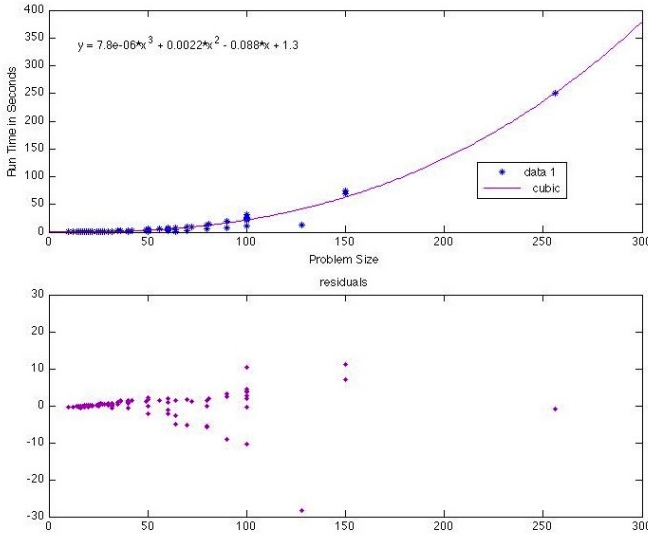

Fig. 2. Performance of QAP as function of number of vertices. The top panel depicts run time, and the bottom panel depicts residual errors. Each point represents a single simulation. Data was sampled from the following model: XXX conroy: you know the answer to this? XXX.

3.4 Brain-Graph Matching

A "connectome" is a brain-graph in which vertices correspond to (collections of) neurons, and edges correspond to connections between them. The *Caenorhabditis elegans*

TABLE 1
Brain-graph matching summary statistics. The maximum number of restarts for both was 30. That the median number of restarts for the electrical connectome is less than 30 implies that our approach found the optimal solution many times, just not always.

	chemical	electrical
Median Accuracy	100%	59.5%
Median # Restarts	3	26
Avg. Solution Time	182 sec.	226 sec.

(*C. elegans*) is a small worm (nematode) with 302 labeled vertices. We consider the subgraph with 279 somatic neurons that form edges with other neurons [?], [6]. Two distinct kinds of edges exist between vertices: chemical and electrical "synapses" (edges). Any pair of vertices may have several edges of each type. Moreover, some of the synapses are polyadic, meaning between > 2 vertices. Thus, the connectome of a *C. elegans* may be thought of as a weighted multi-hypergraph, where the weights are the number of edges of each type. The $\mathtt{QAP}_m$ algorithm natively operates on weighted or unweighted graphs. We therefore conducted the following synthetic experiments. Let $A_{uvz} \in \{0, 1, 2, \ldots\}$ be number of synapses from neuron v to neuron u of type z (either chemical or electrical), and let $A_z = \{A_{uv}\}_{u,v \in [279]}$ for $z \in \{e, c\}$ corresponding to the electrical or chemical connectome. Let $B_{iz} = Q_{iz}A_zQ_{iz}^{\mathsf{T}}$, for some Q_{iz} chosen uniformly at random from $\mathcal{Q}$ for $i = 1, \ldots, s$. For each i and z, obtain $\hat{Q}_{iz}$ using $\mathtt{QAP}_m$ as described above. Define "accuracy" as $\frac{1}{279} \sum_{uv} Q_{iz}\hat{Q}_{iz}$. Table 1 shows some summary results of applying $\mathtt{QAP}_m$ to both A_c and A_e for $s = 10$ times. Note that average solution time is actually smaller than predicted via simulations. Further note that while the electrical connectome was more difficult, that the median number of restarts was less than 30. Our stopping criteria on the number of restarts was either (i) perfect assignment or (ii) 30 restarts. Therefore, this approach achieve perfect assignment sometimes even on this harder assignment problem.

To be consistent with previous QAP library benchmarks, we repeated the above analysis using binarized symmeterized versions of the graphs ($A_{uvz} = 1$ if and only if $A_{uvz} \geq 1$ or $A_{vuz} \geq 1$). The resulting summary statistics are nearly identical to those presented in Table 1 (XXX conroy: but time was longer? XXX).

4 DISCUSSION

This work presents an inexact graph matching algorithm based on the Frank-Wolfe algorithm. While others have incorporated the FW algorithm as a subroutine of a graph matching strategy [18], we modified the FW algorithm for GM in a few ways. First, after FW is over, we project the resulting doubly stochastic matrix onto the set of permutation matrices. Second, we initialize the algorithm using either the identity matrix or the doubly flat matrix (the matrix where all elements are

$1/n$). These choices seem to us to be the most obvious places to start if one has to choose. Third, if one of those choices does not work, we restart FW with other "nearby" initial points. These modifications facilitate improved performance on *all* the benchmarks we considered. Moreover, although the algorithm scales cubically with the number of vertices, the leading constants are very small ($\mathcal{O}(10^{-6})$), so the algorithm runs quite fast on reasonably sized networks (e.g., $n \approx 100$). Indeed, on a biologically inspired GM problem, *C. elegans* connectome mapping, this approach was both fast and effective.

Unfortunately, even with very small leading constants for this algorithm, as n increases, the computational burden gets quite high. For example, extrapolating the curve in the top panel of Figure 2, this algorithm would take about 2 years to finish (on a standard laptop from 2009) when $n = 20,000$. We hope to be able to perform GM on graphs much larger than that even, given that the number of neurons in even a fly brain is $\mathcal{O}(10^5)$. Therefore, more efficient implementations are of interest.

A few generalizations of this approach spring to mind. First, that QAP and LAP are so similar suggests that perhaps one could simply implement a single iteration of QAP starting from the identity. While not changing the order of complexity, it could reduce computational time by at least an order of magnitude, without drastically changing performance properties. The relative performance/computational cost trade-off merits further theoretical investigations. Second, the most "costly" subroutine is LAP. Fortunately, LAP is a quadratic optimization problem with linear constraints. A number of parallelized optimization strategies could therefore potentially be brought to bear on this problem [**?**]. Third, for brain-graphs, we have some prior information that could easily be incorporated in the form of vertex attributes. For example, position in the brain, cell type, etc. could be used to measure "dissimilarity" between vertices. The objective function could then be modified:

$$\widetilde{Q}_{AB} = \operatorname*{argmin}_{Q \in \mathcal{D}} \left\| QAQ^{\mathsf{T}} - B \right\|_F^2 + \lambda J(Q) \qquad (16)$$

where $J(Q)$ is a dissimilarity based penalty and λ is a hyper-parameter. Finally, although this approach natively operates on both unweighted and weighted graphs, multi-graphs are a possible extension.

In conclusion, this manuscript has presented a GM algorithm that is fast, effective, and easily generalizable. To facilitate further development and applications, all the code and data used in this manuscript is available from the first author's website, http://jovo.me.

REFERENCES

[1] M. R. Garey and D. S. Johnson, *Computer and Intractability: A Guide to the Theory of NP-Completeness*. W. H. Freeman, 1979.

[2] D. Conte, P. Foggia, C. Sansone, and M. Vento, "THIRTY YEARS OF GRAPH MATCHING," *International Journal of Pattern Recognition and Artificial Intelligence*, vol. 18, no. 3, pp. 265–298, 2004.

[3] J. Richiardi, H. Eryilmaz, S. Schwartz, P. Vuilleumier, and D. Van De Ville, "Decoding brain states from fMRI connectivity graphs." *NeuroImage*, Jun. 2010. [Online]. Available: http://www.ncbi.nlm.nih.gov/pubmed/20541019

[4] L. Wiskott, J.-M. Fellous, N. Kuiger, and C. von der Malsburg, "Face recognition by elastic bunch graph matching," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 19, no. 7, pp. 775–779, Jul. 1997. [Online]. Available: http://ieeexplore.ieee.org/xpl/freeabs_all.jsp?arnumber=598235

[5] O. Sporns, *Networks of the Brain*. The MIT Press, 2010. [Online]. Available: http://www.amazon.com/Networks-Brain-Olaf-Sporns/dp/0262014696

[6] J. White, E. Southgate, J. N. Thomson, and S. Brenner, "The structure of the nervous system of the nematode caenorhabditis elegans." *Philosophical Transactions of Royal Society London. Series B, Biological Sciences*, vol. 314, no. 1165, pp. 1–340, 1986.

[7] B. B. Biswal, M. Mennes, X.-N. Zuo, S. Gohel, C. Kelly, S. M. Smith, C. F. Beckmann, J. S. Adelstein, R. L. Buckner, S. Colcombe, A.-M. Dogonowski, M. Ernst, D. Fair, M. Hampson, M. J. Hoptman, J. S. Hyde, V. J. Kiviniemi, R. Kötter, S.-J. Li, C.-P. Lin, M. J. Lowe, C. Mackay, D. J. Madden, K. H. Madsen, D. S. Margulies, H. S. Mayberg, K. McMahon, C. S. Monk, S. H. Mostofsky, B. J. Nagel, J. J. Pekar, S. J. Peltier, S. E. Petersen, V. Riedl, S. a. R. B. Rombouts, B. Rypma, B. L. Schlaggar, S. Schmidt, R. D. Seidler, G. J. Siegle, C. Sorg, G.-J. Teng, J. Veijola, A. Villringer, M. Walter, L. Wang, X.-C. Weng, S. Whitfield-Gabrieli, P. Williamson, C. Windischberger, Y.-F. Zang, H.-Y. Zhang, F. X. Castellanos, and M. P. Milham, "Toward discovery science of human brain function." *Proceedings of the National Academy of Sciences of the United States of America*, vol. 107, no. 10, pp. 4734–9, Mar. 2010. [Online]. Available: http://www.pubmedcentral.nih.gov/articlerender.fcgi?artid=2842060\&tool=pmcentrez\&rendertype=abstract

[8] E. T. Bullmore and D. S. Bassett, "Brain Graphs: Graphical Models of the Human Brain Connectome." *Annual review of clinical psychology*, no. November 2010, Apr. 2010. [Online]. Available: http://www.ncbi.nlm.nih.gov/pubmed/21128784

[9] J. T. Vogelstein, W. R. Gray, R. J. Vogelstein, and C. E. Priebe, "Graph Classification using Signal Subgraphs: Applications in Statistical Connectomics," *Submitted for publication*, 2011.

[10] O. L. Mangasarian, *Nonlinear Programming (Classics in Applied Mathematics)*. Society for Industrial Mathematics, 1987. [Online]. Available: http://www.amazon.com/Nonlinear-Programming-Classics-Applied-Mathematics/dp/0898713412

[11] R. E. Burkard, S. E. Karisch, and F. Rendl, "QAPLIB A Quadratic Assignment Problem Library," *Journal of Global Optimization*, vol. 10, no. 4, pp. 391–403, 1997. [Online]. Available: http://www.springerlink.com/content/n5468q3g33184443/fulltext.pdf

[12] "Human Connectome Project - NIH Neuroscience Blueprint." [Online]. Available: http://humanconnectome.org/consortia/

[13] "Open Connectome Project." [Online]. Available: http://openconnectomeproject.org/

[14] R. E. Burkard, M. Dell'Amico, and S. Martello, *Assignment Problems*. SIAM, 2009. [Online]. Available: http://books.google.com/books?id=nHIzbApLOr0C\&pgis=1

[15] M. Frank and P. Wolfe, "An Algorithm for Quadratic Programming," *Naval Research Logistics Quarterly*, vol. 3, no. 1-2, pp. 95–110, 1956. [Online]. Available: http://www.ams.org/mathscinet/search/publications.html?pg1=MR\&s1=MR0089102

[16] S. P. Bradley, A. C. Hax, and T. L. Magnanti, *Applied Mathematical Programming*. Addison-Wesley, 1977. [Online]. Available: http://www.amazon.com/Applied-Mathematical-Programming-Stephen-Bradley/dp/020100464X

[17] R. Sinkhorn, "A relationship between arbitrary positive matrices and doubly stochastic matrices," *The Annals of Mathematical Statistics*, vol. 35, no. 2, pp. 876–879, 1964. [Online]. Available: http://www.jstor.org/stable/2238545

[18] M. Zaslavskiy, F. Bach, and J.-P. Vert, "A path following algorithm for the graph matching problem." *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 31, no. 12, pp. 2227–2242, 2009. [Online]. Available: http://eprints.pascal-network.org/archive/00004435/

[19] C. Schellewald, S. Roth, and C. Schnörr, "Evaluation of Convex Optimization Techniques for the Weighted Graph-Matching Prob-

lem in Computer Vision," in *Proceedings of the 23rd DAGMSymposium on Pattern Recognition*. Springer-Verlag, 2001, pp. 361–368.

Joshua T. Vogelstein Joshua T. Vogelstein is a spritely young man, engulfed in a novel post-buddhist metaphor.

John C. Conroy John C. Conroy seems to basically always be right. He also publishes sometimes.

Lou Podrazik

Steve Kratzer

R. Jacob Vogelstein R. Jacob Vogelstein received the Sc.B. degree in neuroengineering from Brown University, Providence, RI, and the Ph.D. degree in biomedical engineering from the Johns Hopkins University School of Medicine, Baltimore, MD. He currently oversees the Applied Neuroscience programs at the Johns Hopkins University (JHU) Applied Physics Laboratory as an Assistant Program Manager, and has an appointment as an Assistant Research Professor at the JHU Whiting School of Engineerings Department of Electrical and Computer Engineering. He has worked on neuroscience technology for over a decade, focusing primarily on neuromorphic systems and closed-loop brainmachine interfaces. His research has been featured in a number of prominent scientific and engineering journals including the IEEE Transactions on Neural Systems and Rehabilitation Engineering, the IEEE Transactions on Biomedical Circuits and Systems, and the IEEE Transactions on Neural Networks.

Carey E. Priebe Buddha in training.